

Learning and Decision Making

Adapted from 2022 course handouts

Pedro U. Lima

Instituto Superior Técnico/Instituto de Sistemas e Robótica

October 2024

Course handouts


All rights reserved

Markov Chains and Markov Decision Processes

A **Markov Chain** (which by definition satisfies the **Markov Property**) with transition probabilities dependent on actions, and with added rewards, is known as **Markov Decision Process (MDP)**

Example: conditional joint probability of state and reward:

$$\Pr\{x_{t+1} = x', r_{t+1} = r | x_t, u_t, r_t, x_{t-1}, u_{t-1}, \dots, r_1, x_0, u_0\} = \Pr\{x_{t+1} = x', r_{t+1} = r | x_t, u_t\}$$



no rewards included in the diagram

Markov Decision Process (MDP)

Given:

- States x
- Actions u
- Transition probabilities $p(x'|u, x)$
- Reward function $r(x, u)$
- Initial state x_0
- *Discount factor* γ
- Horizon T

Goal:

- Find policy $\pi(x)$ that maximizes the expected sum of discounted rewards accumulated over time


$$\max_{\pi} \mathbb{E} \left(\sum_{\tau=0}^T \gamma^{\tau} r_{\tau} \mid \pi \right)$$

Markov Decision Process (MDP)

$$\begin{aligned} x_t &\in X \\ u_t &\in U(x_t) \\ r_{t+1} &\in \mathfrak{R} \end{aligned}$$

Probabilistic Policy function

$$\pi(x_t, u_t) = P(u = u_t | x = x_t), \forall u_t \in U(x_t)$$



Goal: choose the policy that maximizes the

Value function $V_\pi(x) = \mathbb{E} \left(\sum_{\tau=0}^T \gamma^\tau r_\tau \mid \pi, x_0 = x \right)$

T may go to infinity, as long as $\gamma \neq 1$

- State: complete description of the state of the world.
 - whole chess board information in a chess game.
 - position, velocity, angle, angular velocity of a cart-pole system.
- Action: possible actions in the environment.
 - discrete: left, right, up, down.
 - continuous: robot wheel velocities.
- Reward: $r(x_t, a_t)$ measure how “good” an action for a particular state is.
 - angle of the cart-pole close to zero.
- Discounted cumulative reward:

$$R_t = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots$$

Policy (fully observable case) is a map of states onto actions:

$$\pi : x_t \rightarrow u_t$$

Expected discounted cumulative reward / payoff:

$$R_T = E \left[\sum_{\tau=0}^T \gamma^{\tau} r_{t+\tau+1} \right], \quad 0 < \gamma \leq 1$$

- $T=0$: greedy policy
- $T>0$: finite horizon case, typically no discount
- $T=\infty$: infinite-horizon case, finite reward if discount $\gamma < 1$

- **Goal:** Find policy $\pi(x)$ that maximizes the expected cumulative reward
- Two methods for finding the optimal policy:
 - value iteration
 - policy iteration

Policies cont'd.

- Expected cumulative payoff of policy π , for a given state x :

$$R_T^\pi(x) = E \left[\sum_{t=0}^T \gamma^t r_t \mid x_0 = x, \pi \right]$$

- Optimal value function:

$$V^*(x) = \max_{\pi} R_T^\pi(x) = \max_{\pi} E \left[\sum_{t=0}^T \gamma^t r_t \mid x_0 = x, \pi \right]$$

- Optimal policy:

$$\pi^* = \arg \max_{\pi} R_T^\pi(x)$$

Value Iteration

- 1-step optimal value function and policy

$$V_1^*(x) = \max_u \sum_{x'} P(x'|x, u) r(x, u, x')$$

$$\pi_1^*(x) = \arg \max_u \sum_{x'} P(x'|x, u) r(x, u, x')$$

- 2-step optimal value function and policy:

$$V_2^*(x) = \max_u \sum_{x'} P(x'|x, u) (r(x, u, x') + \gamma V_1^*(x'))$$

$$\pi_2^*(x) = \arg \max_u \sum_{x'} P(x'|x, u) (r(x, u, x') + \gamma V_2^*(x'))$$

k-step Value Iteration

- Optimal Value function:

$$V_k^*(x) = \max_u \sum_{x'} P(x'|x, u) (r(x, u, x') + \gamma V_{k-1}^*(x'))$$

- Optimal Policy:

$$\pi_k^*(x) = \arg \max_u \sum_{x'} P(x'|x, u) (r(x, u, x') + \gamma V_{k-1}^*(x'))$$

- Optimal policy, infinite horizon:

$$V_{\infty}^*(x) = \max_u \sum_{x'} P(x'|x, u) (r(x, u, x') + \gamma V_{\infty}^*(x'))$$

- Bellman equation
- Fixed point is optimal policy
- Necessary and sufficient condition:
induced policy is optimal iff
value function satisfies the above condition

Algorithm

For all x

$$V_0^*(x) = 0$$

For $k = 1, 2, \dots$ until convergence

For all x

$$V_k^*(x) = \max_u \sum_{x'} P(x'|x, u) (r(x, u, x') + \gamma V_{k-1}^*(x'))$$

Value iteration converges to the optimal value function, which satisfies the Bellman equation.



$$V_k^* \rightarrow V^*(x)$$

$$V^*(x) = \max_u \sum_{x'} P(x'|x, u) (r(x, u, x') + \gamma V^*(x'))$$

The optimal policy is given by

$$\pi^*(x) = \arg \max_u \sum_{x'} P(x'|x, u) (r(x, u, x') + \gamma V^*(x'))$$

Value Iteration for Motion Planning



Value Function and Policy Iteration

- Often the optimal policy has been reached long before the value function has converged.
- Policy iteration calculates a new policy based on the current value function and then calculates a new value function based on this policy.
- This process often converges faster to the optimal policy.

Policy Iteration

repeat

repeat (keeping the same policy)

for all x

$$V_{i+1}^{\pi_k}(x) = \sum_{x'} P(x'|x, \pi_k(x)) (r(x, \pi_k(x), x') + \gamma V_i^{\pi_k}(x'))$$

policy evaluation

until value function has converged

update policy

for all x

$$\pi_{k+1}(x) = \arg \max_u \sum_{x'} P(x'|x, u) (r(x, u, x') + \gamma V^{\pi_k}(x'))$$

policy improvement

until policy has converged

Previous (DP) methods to solve MDPs assume full knowledge of $p(x'|u,x)$ and $r(u,x)$

Dynamic Programming (DP)

- To determine V for $|X| = N$, a system of N non-linear equations must be solved.
- Well-established mathematical method.
- A complete model of the environment is required (P and R known).
- Often faces the “*curse of dimensionality*” [Bellman, 1957]

Alternative approaches, if we do not know $p(x'|u,x)$ and $r(u,x)$

Monte Carlo

- Similar to DP, but P and R_s unknown.
- P and R determined from the average of several trial-and-error trials.
- Unappropriate for a step-by-step incremental approximation of V^* .

Temporal Differences

- Knowledge of P e R is not required
- Step-by-step incremental approximation of V .
- Mathematical analysis more complex.
- examples: *Q-learning*, *SARSA*, ...

Should one learn $V^*(x)$?

- The agent should prefer a state with higher V , because the future cumulative reward will be greater
- But the agent chooses actions, not states
- All fine, then

$$\pi^*(x) = \operatorname{argmax}_u \{ r(x, u) + E[\gamma V^*(\delta(x, u))] \}$$

Unknown!

Value Functions

state value for policy π :

$$V_{\infty}^{\pi}(x) = E_{\pi} \left\{ \sum_{k=0}^{\infty} \gamma^k r_{t+k+1} \mid x_t = x \right\}$$

Expected value of starting in state x and following policy π thereafter.

NOTE: value of final state, if any, is always zero.

(state, action) value for policy π :

$$Q_{\infty}^{\pi}(x, u) = E_{\pi} \left\{ \sum_{k=0}^{\infty} \gamma^k r_{t+k+1} \mid x_t = x, u_t = u \right\}$$

Expected value of starting in state x , *carrying out action* u , and following policy π thereafter.

Learn $Q^*(x, u)$ instead!

Value Functions cont'd

relation between state value and Q function for policy π :

Q is such that its value is the maximum discounted cumulative reward that can be achieved starting from state x and applying action u as the first action

$$Q(x, u) = E[r_{t+1} + \gamma V^*(x_{t+1}) | x_t = x, u_t = u]$$

$$\pi^*(x) = \operatorname{argmax}_u Q(x, u)$$

$$V^*(x) = \max_u Q(x, u)$$

$$\therefore Q(x, u) = E[r_{t+1} + \gamma \max_{u'} Q(x', u') | x_t = x, u_t = u]$$

Value Functions cont'd

Bellman equation for V and Q
(discrete action and state spaces, deterministic policy)

$$V^*(x) = \max_u \sum_{x'} P(x'|x, u) (r(x, u, x') + \gamma V^*(x'))$$

$$Q^*(x, u) = \sum_{x'} P(x'|x, u) (r(x, u, x') + \max_{u'} \gamma Q^*(x', u'))$$

Value iteration

$$V_k^*(x) = \max_u \sum_{x'} P(x'|x, u) (r(x, u, x') + \gamma V_{k-1}^*(x'))$$

Q-value iteration

$$Q_{k+1}(x, u) = \sum_{x'} P(x'|x, u) (r(x, u, x') + \max_{u'} \gamma Q_k(x', u'))$$

Value Functions cont'd

Q-value iteration

$$Q_{k+1}(x, u) = \sum_{x'} P(x'|x, u) (r(x, u, x') + \max_{u'} \gamma Q_k(x', u'))$$

Rewritten as an expectation

$$Q_{k+1}(x, u) = E_{x' \sim P(x'|x, u)} [r(x, u, x') + \max_{u'} \gamma Q_k(x', u')]$$

Tabular Q-learning:

- Replace expectation by samples

$$x' \sim P(x'|x, u)$$

- Compute error w.r.t Bellman equation and iterate Q-value

$$Q_{k+1}(x, u) = Q_k(x, u) + \alpha \left[r(x, u, x') + \max_{u'} \gamma Q_k(x', u') - Q_k(x, u) \right]$$

└─ Learning rate

Tabular Q-Learning - Algorithm

Algorithm:

Initialize $Q_0(x, u)$ for all x and u

Initialize current state x

For $k=1, 2, \dots$ until convergence

 Sample action u

 Execute action u , get r and x'

 Compute $Q_{k+1}(x, u)$

$$Q_{k+1}(x, u) = Q_k(x, u) +$$

$$\alpha \left[r(x, u, x') + \max_{u'} \gamma Q_k(x', u') - Q_k(x, u) \right]$$

 Update current state $x \leftarrow x'$

α_n constant allows adaptability to slow environment changes but it does not guarantee convergence – only possible with a temporal decay under given circumstances.

Algorithm Convergence

- All states and actions are visited infinitely often
- Learning rate is such that

$$0 < \alpha_k < 1$$

$$\sum_{k=0}^{\infty} \alpha_k(x, u) = \infty$$

$$\sum_{k=0}^{\infty} \alpha_k(x, u)^2 < \infty$$

Then $\forall x, u \ P[\lim_{k \rightarrow \infty} Q_k(x, u) = Q^*(x, u)] = 1$

Action Selection:

Exploration vs Exploitation

Exploration: less promising actions, which may lead to good results, are tested.


Exploitation: takes advantage of tested actions which are more promising, i.e., which have a larger $Q(x,u)$.

- *ϵ - greedy:* at each step n , picks the best action so far with probability $1-\epsilon$, for small ϵ , but can also pick with probability ϵ , in an uniformly distributed random fashion, one of the other actions.


- *softmax:* at each step n , picks the action to be executed according to a Gibbs or Boltzmann distribution:

$$\pi_n(x,u) = \frac{e^{Q_n(x,u)/\tau}}{\sum_{u'(x)} e^{Q_n(x,u')/\tau}}$$

Q-Learning – an Example




$r(x,u)$
 $V^*(x)$




$Q_n^\pi(x,u)$
 $\alpha_n = 1$
 $\gamma = 0.9$

Q-Learning – an Example



$r(x,u)$
 $V^*(x)$

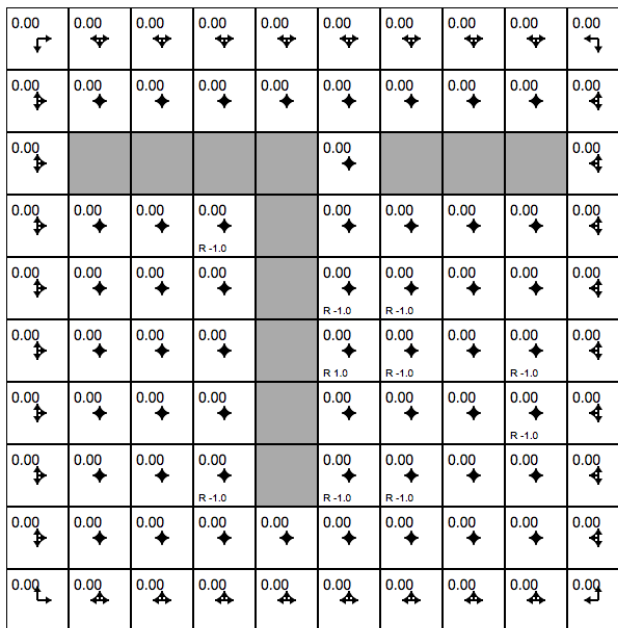


$Q_n^\pi(x,u)$

Q-Learning – Grid World

<http://cs.stanford.edu/people/karpathy/reinforcejs/index.html>


Exploration epsilon: 0.20



Setup



Dynamic Programming



Q-Learning

Tabular methods do not scale

Discrete environments (number of states)

- Tetris: 10^{60}
- Atari games: 10^{308}

Discretized continuous environments

- Inverted pendulum: 10^2
- Hopper: 10^{10}
- Humanoid: 10^{100}


Alternative to tabular representation:

- Parametrized Q function $Q_{\theta}(x, u)$
Typically a neural network.

From tabular Q-learning to DQN

Alternative to tabular representation:

- Parametrized Q function $Q_{\theta}(x, u)$
Typically a neural network.



- New update rule on the parameters (gradient-based)

$$\theta_{k+1} = \theta_k - \alpha \Delta_{\theta} \left[\frac{1}{2} (Q_{\theta_k}(x, u) - r(x, u, x') - \gamma \max_{u'} Q_{\theta_k}(x', u'))^2 \right]$$

Alternative to sampling once and updating the Q-function

- Sample and store several actions in a replay memory
- Select a small batch from the replay memory and perform gradient descent using that batch

Many methods out there

- Deep Q-learning (DQN)
 - Learn the Q function (parametrized by θ)
 - Policy π is generated directly from Q
- Policy gradient methods
 - Directly learn π (parametrized by θ)
 - TRPO (Trust Region Policy Optimization)
 - PPO (Proximal Policy Optimization)
- Actor Critic methods
 - Neural nets for the value function and the policy
 - DDPG (Deep Deterministic Policy Gradient)
 - SAC (Soft Actor Critic)

Real Robot RL (Q-learning)



Real Robot System (MDP + RL)



References:

- Sutton, Richard S., and Andrew G. Barto. Introduction to reinforcement learning. Vol. 135. Cambridge: MIT Press, 1998.
- Mitchell, Thomas M. "Machine Learning." (1997).
- Sebastian Thrun, Wolfram Burgard and Dieter Fox, *Probabilistic Robotics*, 2005 The MIT Press
- M. T. J. Spaan, "Partially Observable Markov Decision Processes", in Reinforcement Learning: State of the Art, M. A. Wiering and M. van Otterlo, editors, Springer Verlag, 2012.